Pranav Dhiran

pranavdhiran.me · linkedin.com/in/prannav-dhiran · github.com/Pranav-d33 · Nagpur, Maharashtra · +91-8999629839 · dhiranpranav72@gmail.com

PROFILE

Final-year ECE undergrad with hands-on experience building LLM-based systems — transformer pre-training to multi-agent pipelines and open-source LLM-hardware integrations. LFX '26 Mentee (Hyperledger Cello) and active CNCF contributor; looking to work on large, production-grade codebases where there is real engineering to learn from.

EDUCATION

B.Tech — Electronics & Telecommunication Engineering 2023 – Present
SGGS Institute of Engineering & Technology, Nanded | Minor in IT | Final Year

EXPERIENCE

LFX Mentee | *Hyperledger Cello · Linux Foundation* Jun 2026 – Present

- Hyperledger Cello manages the full lifecycle of Hyperledger Fabric networks — every operation today requires manual dashboard interaction, creating a steep learning curve.
- Building an AI agent that takes natural language input, reasons over the right sequence of Cello API calls, and executes them — turning network operations and Chaincode debugging into a conversation.

Research Intern | *IRT, University of South Carolina* Apr 2026 – Present

- Researching neurosymbolic SLM architecture and pre-training pipelines — integrating symbolic reasoning constraints into small language model training for structured knowledge tasks.
- Working on RL-based fine-tuning (GRPO/RLHF) for SLM alignment — reward modeling, policy optimization, and evaluation on neurosymbolic reasoning benchmarks.

Open Source Contributor | *Meshery — CNCF Sandbox Project* Mar 2026 – Present

- 5+ merged PRs — service mesh management features, UI components, and API integrations across Go backend and React frontend.
- Active in code reviews, issue triage, and community discussions per CNCF contributor guidelines.

PROJECTS

Small Language Model from Scratch — *PyTorch · Hugging Face · AMP · NLP* · GitHub

- Pre-trained a GPT-style autoregressive transformer (6L, 6H, 384-dim) on TinyStories — custom BPE tokenization, mmap pipelines, AMP mixed precision, gradient accumulation, cosine LR; systematic depth/embedding ablations for scaling behavior analysis.

Medaura — Agentic Pharmacy System — *FastAPI · LangChain · ChromaDB · Langfuse · React* · Live

- 5-agent autonomous pharmacy system (Ordering, Safety, Forecast, Procurement, UI) with semantic RAG via ChromaDB + sentence-transformers and Langfuse observability — live on web, zero human intervention, 4 languages.

GNU Radio MCP Server — LLM-to-SDR Bridge — *Python · FastMCP · ZMQ · XML-RPC · GNU Radio* · GitHub

- Provider-agnostic MCP adapter bridging LLMs to live GNU Radio SDR flowgraphs — 13 tools, Pydantic v2 validation, ZMQ context management, async IQ capture with Welch PSD; natural-language control of HackRF/RTL-SDR hardware.

RF Watch — Open-Source RF Spectrum Monitor — *Python · GNU Radio · HackRF One · Signal Processing* · GitHub

- Real-time RF spectrum analyzer — FFT-based spectral feature extraction with lightweight ML classifier for passive transmitter detection; born from SIH 2025 anti-drone work for ITBP.

TECHNICAL SKILLS

Languages & Frameworks: Python, Go (learning), PyTorch, TensorFlow, Hugging Face Transformers, TRL

APIs & Protocols: REST, gRPC, GraphQL, OpenAI/Gemini/Ollama APIs, FastMCP, XML-RPC, ZMQ

LLM Engineering: Instruction Fine-Tuning, LoRA, PEFT, GRPO, RLHF, INT4/INT8 Quantization, Unsloth, LangChain, LangGraph, ChromaDB, FAISS

Agentic & Infra: Multi-Agent Systems, Tool-Use, Function Calling, MCP Servers, LangSmith, Langfuse, Docker (basic), Kubernetes (basic)

CERTIFICATIONS

Fine-tuning & RL for LLMs: Intro to Post-training — DeepLearning.AI · AI Agents in LangGraph — DeepLearning.AI · Quantization Fundamentals — Hugging Face · MCP: Build Rich-Context AI Apps — Anthropic

AWARDS & RECOGNITION

- International Finalist (Top 6) — UWA Hack for Impact 2026
- National Finalist — Smart India Hackathon (SIH) 2024 & 2025
- Regional Qualifier — Nxt Wave × OpenAI Buildathon